


Van Hateren null rotation across feature dimension ($n = 1000$, $\sigma^2/S = 0.01$)



The $p=16$ metric is unchanged from the original experiment; larger p appends disjoint prediction-null features. MC uses the fixed DE-CV ridge penalty.